

01-working-demo-runbook

S3: Working Demo

Revix · *driven by reviews*. The runbook: what to show, in what order, on what, and what to do when something breaks.

Submission	S3 Working demo
Weight	20 marks of 100
Due	Friday 16 October 2026
Marked on	The application should work. A video may serve as backup
Team	Aditya Nariyapara, Devika Jonjale, Saachi Shinde
Live site	<code>https://revix-reviews.vercel.app</code>
Live API	<code>https://revix-api-tcyq.onrender.com</code>
Companion	Speaking script

Read the rubric literally. The mark is for the application *working*, not for novelty. The demo is therefore built around one uninterrupted path that is guaranteed to work, with the interesting material offered as detours we can take if there is time. Never open with something that might fail.

1. The single most important decision

Demo the deployed site, not localhost. A live URL the examiner can open on his own phone is worth more than anything running on your laptop, and it removes "does it only work on your machine" as a question entirely.

Record the backup video anyway. The rubric explicitly permits it, which means it is expected, and it costs one afternoon. See [§7](#).

2. Pre-flight, the morning of

Run through this list. It takes four minutes and it is the difference between a demo and an incident.

#	Check	Command or action	Expected
1	The API is awake and healthy	<code>curl -s \$API/health</code>	<code>"status":"ok"</code> , <code>"database":true</code> , non-zero counts
2	All three sources are healthy	<code>curl -s \$API/sources/health</code>	No source reporting <code>circuit_open</code>
3	The nightly run succeeded	GitHub Actions, latest nightly	Green
4	<code>main</code> is green	GitHub, CI checks	All six jobs pass
5	The site is warm	Load the home page twice	Second load is instant
6	Your demo variant still resolves	Open its verdict page directly	Loads, unsuppressed, has evidence
7	The comparison you plan to show still exists	Open <code>/compare</code> and click the pair	Both sides populated
8	Phone check	Open the site on a phone	Layout holds

Warm the site fifteen minutes before, and again five minutes before. The free tier idles the instance out. This single habit prevents the most likely failure of the day.

Pick your demo variant the night before, not live. Choose one with a high evidence count and visible disagreement, since disagreement is what makes the product's argument. Write the URL here before the demo:

Demo variant URL: _____

Backup variant URL: _____

3. The demo flow

Ten minutes of screen time. Each step names what to click and what the examiner should notice.

Step 1: The home page (30 seconds)

Show: `https://revix-reviews.vercel.app`

State the problem in one sentence, then move. Do not linger on the landing page; nothing is being marked here.

Step 2: Browse the catalogue (1 minute)

Click: *Browse*

What to point out: the scale is real. **143 variants across 42 models and 16 manufacturers**, cars and two-wheelers, every one of them carrying a verdict built from real reviews.

Show the filters working: vehicle class, and the free-text search. Type a model name and let the list narrow. This proves the catalogue is queryable, not a fixed list.

Step 3: A verdict (2 minutes), the core of the demo

Click: into your chosen variant.

Three things to point at, in this order:

1. **The score has a range, not just a number.** That is a bootstrap confidence interval, and it is on screen because a single decimal would be a lie about precision.
2. **The evidence count and the effective sample size.** These are different numbers on purpose: the effective sample size accounts for how unevenly the reviews are distributed, so a verdict built from one very loud source is visibly weaker than one built from many.
3. **The aspects, ordered with the most-disagreed-upon at the top.** Disagreement is the useful signal, so it is what the layout leads with.

Step 4: Traceability (2 minutes), the differentiator

Click: any score, to open its evidence.

This is the claim the whole project rests on. Every number opens the actual sentences it was computed from, with the source, the link and the weight each one contributed.

Point out that the citations are ranked by contribution weight, and that these rows were written by the fusion engine *before* any prose existed. The score was computed *from* them, which is why the citation cannot be wrong. It is not a language model asked to justify a number after the fact.

Follow one link out to the original review on CarWale or CarDekho, so he sees a real page on a real site. Doing this once is worth a paragraph of explanation.

Step 5: The weighting switch (1.5 minutes), the "wow" moment

Back on the verdict page, flip the weighting between equal, source-weighted and credibility-weighted.

Watch the numbers move. That is the same evidence under three different trust models.

Then the line that matters:

This is instant because nothing is computed here. All three verdicts were calculated by the pipeline overnight and stored. Switching is a lookup, not a recomputation. That is the same architectural decision that gives us our latency and our demo reliability.

Step 6: Comparison (1 minute)

Click: *Compare*, pick a suggested pair.

Pairs are generated by taking one variant per model and pairing price neighbours within a vehicle class, so the suggestions are vehicles a person would actually cross-shop.

Step 7: Honesty pages (1.5 minutes)

These carry more marks than they look like they do, because they answer "how do we know any of this is true."

Sources: where every source stands, with unit counts and last success. A dead source degrades the system; it does not break it.

Status: the pipeline's own health.

Accuracy: our published measurements, *including the unflattering one*. The aspect classifier we built lost to the lexicon it was trained from and ships disabled. **Show this deliberately.** A project that publishes the result that went against it is making a stronger claim than one that only publishes wins.

Method: how a verdict is built, in plain language.

Step 8: Suppression (30 seconds)

Show a variant held back below the evidence floor.

Nine of our 143 variants do not have enough evidence to say anything responsibly, so we say nothing and explain why, rather than showing a confident-looking number built from four reviews.

Refusing to answer is a feature. Make sure it is seen.

4. Optional detours, only if invited

Do not volunteer these. Have them ready if he asks for more.

If he asks about	Show
Speed	Browser dev tools, network tab: the <code>X-Response-Time-ms</code> header on a response
Security	The response headers: CSP, HSTS, nosniff
Failure behaviour	<code>/health</code> and <code>/sources/health</code> returning real JSON
The API itself	<code>/docs</code> , the generated OpenAPI page
Accessibility	Tab through the verdict page with the keyboard only
Mobile	Hand him your phone with the site already open

5. What can break, and what to say

Failure	Say this	Then do this
Site is slow to first load	"That is the free-tier cold start; the instance idles out and we ping it on a schedule to reduce it"	Keep talking, it wakes
A page errors	"That is our error boundary rather than a blank page"	Reload; move to the next step
API is down entirely	"This is why we brought the video"	Switch to the recording without ceremony
A source shows <code>circuit_open</code>	"A source failed and the site kept serving; the status page shows it"	Open <code>/sources</code> and turn it into the reliability point
Data looks different from this document	"The pipeline ran last night, so the numbers move"	Read the live numbers off the screen
Something genuinely is broken	"That is a real bug, and here is what I think it is"	Diagnose honestly. Do not pretend it is intended

The rule: never say "it usually works." Either show it, or say plainly that it is broken and what you would check first. A calm diagnosis reads better than a flustered excuse, and this examiner will have seen a hundred flustered excuses.

6. Equipment

Laptop, charged, charger present	
Browser: one window, demo tabs only	Close everything else, including notifications
Zoom level ~110% so text is readable on a projector	
Phone with the site loaded, for the responsive check	
Backup video on the laptop and on a phone or drive	Not only in the cloud
This runbook, printed or on a second device	Not on the screen you are sharing
Hotspot ready, in case the room's network fails	

7. The backup video

Record it **a week before**, not the night before, so there is time to redo it.

Length	5 to 6 minutes
Content	Steps 1 to 8 of §3, in order, no detours
Audio	Narrated. A silent screen recording does not carry the argument
Quality	1080p, cursor visible, no notifications on screen
Store	On the laptop, on a drive, and uploaded. Test playback on the laptop you will bring
Name	<code>revix-working-demo-<date>.mp4</code>

Record it in one take if you can. A video with one small stumble looks like software; a heavily edited one looks like a trailer.

8. Rehearsal

Rehearse the full flow **at least twice**, once end to end without stopping, and once with someone interrupting to ask questions. Time it. If it runs past twelve minutes, cut step 6 rather than rushing step 4, because traceability is the mark and comparison is a nicety.

Rehearse **on the room's setup if you can get in beforehand**, or at least on an external display, since projector colour and text size are the two things that surprise people.

9. Refreshing the numbers in this document

```
API=https://revix-api-tcyq.onrender.com
curl -s $API/health
curl -s $API/sources/health
curl -s "$API/variants?limit=200" | python -c "import sys,json; d=json.load(sys.stdin);
print(len(d),'variants,',sum(1 for v in d if not v['is_suppressed']),'published')"
```

Figures quoted here were taken on **7 September 2026** and will have grown. Read the live ones off the screen during the demo rather than quoting this page.