

02-working-demo-speaking-script

S3 Working Demo: Speaking Script

For the one-to-one demo with Prof. Naik. Companion to [the runbook](#).

Covers	S3 Working demo, 20 marks
Marked on	The application should work
Format	One-to-one, screen shared, live site
Target length	10 minutes of demo, then questions

How to use this

Each section gives the click, then the words. **Bold sentences carry the mark**, so if you are running short, say only those.

Talk while you click. Silence while a page loads reads as waiting for something to go wrong. Narration reads as confidence.

Say what he should notice before he notices it. An examiner looking at the wrong part of the screen is a mark lost to nothing.

0. Opening (30 seconds)

Site already loaded and warm. Do not start with a blank browser.

Revix is live at this URL and you can open it on your own phone afterwards.

The problem is that a star average hides who rated, what they cared about, and whether they agreed. **Revix rebuilds the verdict from the reviews underneath it, and every single number on this site opens the sentences it was computed from.**

It is running on **real data collected from three sources**, refreshed by a pipeline that runs itself every night. I will take about ten minutes to walk the main path.

1. Browse (1 minute)

Click: *Browse*

This is the catalogue: **143 variants across 42 models and 16 manufacturers**, cars and two-wheelers.

Use the filter, then type in the search.

It filters by vehicle class and searches across manufacturer, model and variant. **Every one of these carries a verdict built from real owner reviews**, not seeded or sample data.

2. A verdict (2 minutes), *slow down here*

Click: into your chosen variant.

This is the core screen. Three things I would point at.

Point at the range.

The score has a range, not just a number. That is a bootstrap confidence interval. A single decimal would claim a precision we do not have, so we show the interval instead.

Point at the evidence counts.

These two numbers are different on purpose. The evidence count is how many review units went in. **The effective sample size accounts for how unevenly they are distributed**, so a verdict built mostly from one very loud source is visibly weaker than one built evenly from many. You can see the quality of the verdict, not just the verdict.

Point at the aspect ordering.

And the topics are ordered with **the most disagreed-upon at the top**, because disagreement is the useful signal. If owners all agree the mileage is fine, that tells you less than the fact that they are split on the gearbox.

3. Traceability (2 minutes), the mark

Click: any score, to open the evidence.

This is the claim the whole project rests on. Every number opens the actual sentences it was computed from: the source, the link, and how much each one contributed.

Point at the ordering.

They are ranked by contribution weight. **And these rows were written by the fusion engine before any text existed**, and the score was computed *from* them. That is why the citation cannot be wrong.

It is not a language model asked to justify a number after the fact. The evidence is the input, not an explanation generated afterwards.

Click through to the original review on the source site.

And that is the real review, on CarWale, written by a real owner.

Doing this once is worth a paragraph of explanation. Do not skip it.

4. The weighting switch (1.5 minutes), *the moment they remember*

Back on the verdict, flip between the three weightings.

Now watch the numbers. **That is the same evidence under three different trust models:** everything counted equally, weighted by source, and weighted by reviewer credibility.

Pause for a second and let him see it move.

And notice it is instant, because nothing is being computed here. All three were calculated by the pipeline overnight and stored as finished rows. Switching is a lookup.

That is the same decision that makes this demo safe to give. No model runs during a request, so this page cannot fail because somebody else's API is down or over quota.

5. Compare (1 minute)

Click: *Compare*, then a suggested pair.

The suggestions are generated by taking one variant per model and pairing price neighbours within a class, so these are vehicles somebody would actually cross-shop rather than arbitrary pairs.

6. The honesty pages (2 minutes), *worth more marks than they look*

Click: *Sources*

Where every source stands: unit counts, last successful run, error rate. **A source that fails degrades the system, it does not break it:** the other sources carry on and this page shows you the gap rather than hiding it.

Click: *Accuracy*

These are our published measurements, and I want to show you one specifically.

We built a machine-learning aspect classifier and it lost to the simple lexicon it was trained from, by 0.29 macro F1. So it ships disabled, and the result is published here.

We kept the code and the number rather than quietly dropping it. **A project that publishes the result that went against it is making a stronger claim than one that only publishes its wins**, and the whole argument of Revix is that a number should be earned.

Click: *Method*

And this explains how a verdict is built, in plain language, for a reader who is not going to read our code.

7. Refusing to answer (30 seconds)

Show a suppressed variant.

Last thing. **Nine of our variants do not have enough evidence to say anything responsibly, so we say nothing** and explain why, instead of showing a confident-looking score built from four reviews.

Refusing to answer is a feature, not a gap in the data.

8. Closing (20 seconds)

So: live, on real data, refreshed nightly without anybody touching it, with every number traceable to its evidence and the uncertainty on screen rather than hidden.

Happy to go into any part of it, or you can open it on your phone now.

Questions to expect

"Is this real data or seeded?"

Entirely real. Owner reviews from CarWale and BikeWale, CarDekho and BikeDekho, and comments on YouTube review videos, all collected under each source's own rules. The counts on the sources page are live.

"What happens if a source goes down?"

The pipeline records the failure and carries on with the others, and the sources page shows it. That is what the circuit breaker is for. Nothing about the site's availability depends on a source being up, because ingestion happens overnight and the site only reads finished rows.

If a source is actually showing a failure on the day, open it and use it. A demonstrated graceful degradation is better than a described one.

"Only three sources?"

Three that permit us. Team-BHP returns 403 to our agent and two other sites disallow us outright in robots.txt. **We chose to respect that rather than spoof a browser.** Our connectors identify honestly and rate-limit themselves. Working around it would mean misrepresenting who we are.

"How do I know the evidence actually produced that score?"

Because the citation rows are written by the fusion engine before any prose exists, and the score is computed from those rows. The evidence is the input. There is no step where a model is asked to explain a number it did not produce.

"Why is that variant empty?"

It is below our evidence floor: fewer than forty evidence units, or fewer than two independent sources. We hold it back rather than publish something we cannot stand behind, and the page says which condition it failed.

"Is it fast because the dataset is small?"

No, and that is the architecture rather than the size. Nothing is computed during a request; every endpoint reads a row the pipeline already wrote. **Ten times the data would change how long the night takes and change nothing about this page.**

"Can I try it myself?"

Please do. It works on a phone and it does not need a login.

This is the best question you can get. Hand him the phone.

Recovery lines

Keep these ready. Delivered calmly, a failure costs almost nothing; delivered badly, it costs the room.

Situation	Line
First load is slow	"That is the free-tier cold start; the instance idles out and we ping it on a schedule to reduce it."
A page errors	"That is our error boundary catching it rather than showing a blank page."
Numbers differ from the runbook	"The pipeline ran last night, so these move. Let me read the live ones."
A source is failing	"Good timing: this is exactly the degradation I described. The site is still serving, and the status page shows the gap."
Everything is down	"This is why we recorded a backup." <i>Switch without apologising twice.</i>
A genuine bug appears	"That is a real bug. My first guess is X, and I would check Y."

Apologise once, briefly, then keep moving. Repeated apologising makes a small problem look like a large one.